

# Neural Networks

Prof. Sharon McNicholas

STATS 780/CSE 780

1

## Introduction

- As with many other topics covered in an hour or two during this course, there are whole courses and entire books devoted to neural networks.
- Like PPR, a neural network is a non-linear statistical model.
- Originally developed as a model of the brain.
- A neural network is a two-stage model that can be used for regression or classification; again, we want to use  $\mathbf{X}$  to predict  $Y$ .


2

## The Idea

- Similar to Hastie et al. (2009)<sup>a</sup>, we will consider classification using a neural network with one hidden layer (aka a single-layer perceptron).
- Suppose  $Y$  has  $K$  classes and take  $\mathbf{X} = (X_1, \dots, X_p)'$ .
- For our purposes, consider binary variables  $Y_1, \dots, Y_K$ .
- Introduce hidden units  $Z_1, \dots, Z_M$ , where each  $Z_m$  is a linear combination of  $X_1, \dots, X_p$ .
- Here is a visual representation of a neural network from Hastie et al. (2009); note that bias can be added to each unit in the hidden and output layers.

<sup>a</sup>Hastie, T., Tibshirani, R. and Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Second Edition. Springer: NY.

## Example of Neural Network



**FIGURE 11.2.** Schematic of a single hidden layer, feed-forward neural network.

## The Idea contd.

- Similar to Hastie et al. (2009), we can describe the neural network as follows:

$$Z_m = \sigma(\alpha_{0m} + \boldsymbol{\alpha}'_m \mathbf{X}),$$

$$T_k = \beta_{0k} + \boldsymbol{\beta}'_k \mathbf{Z},$$

$$f_k(\mathbf{X}) = g_k(\mathbf{T}),$$

for  $m = 1, \dots, M$  and  $k = 1, \dots, K$ , where  $\mathbf{Z} = (Z_1, \dots, Z_M)'$  and  $\mathbf{T} = (T_1, \dots, T_K)'$ .

- The  $\alpha_{0m}$ ,  $\boldsymbol{\alpha}_m$ ,  $\beta_{0k}$ , and  $\boldsymbol{\beta}_k$  are unknown parameters called weights.
- $\sigma(\cdot)$  is a (non-linear) transformation known as the activation function.
- $g_k(\mathbf{T})$  allows a transformation of  $\mathbf{T}$  and is called the output function.

5

## Activation Function

- Suppose  $\sigma(\cdot)$  is the identity function, then we have

$$Z_m = \alpha_{0m} + \boldsymbol{\alpha}'_m \mathbf{X}$$


and so a linear model.

- For  $\sigma(\cdot)$  non-linear, a neural network can be viewed as a non-linear generalization of a linear model.
- The sigmoid function is a popular choice for  $\sigma(\cdot)$ , i.e.,

$$\sigma(v) = \frac{1}{1 + e^{-v}}.$$

6

## Sigmoid Function (from Hastie et al., 2009)



**FIGURE 11.3.** Plot of the sigmoid function  $\sigma(v) = 1/(1 + \exp(-v))$  (red curve), commonly used in the hidden layer of a neural network. Included are  $\sigma(sv)$  for  $s = \frac{1}{2}$  (blue curve) and  $s = 10$  (purple curve). The scale parameter  $s$  controls the activation rate, and we can see that large  $s$  amounts to a hard activation at  $v = 0$ . Note that  $\sigma(s(v - v_0))$  shifts the activation threshold from 0 to  $v_0$ .

7

## Output Function

- For regression,  $g_k(\mathbf{T})$  is usually chosen to be the identity function, i.e.,

$$g_k(\mathbf{T}) = T_k.$$

- For classification, the softmax function is common:

$$g_k(\mathbf{T}) = \frac{e^{T_k}}{\sum_{h=1}^K e^{T_h}}. \quad (1)$$

- For any  $T_k$ , (1) guarantees that  $g_k(\mathbf{T}) > 0$ , for all  $k = 1, \dots, K$ , and

$$\sum_{k=1}^K g_k(\mathbf{T}) = 1.$$

8

## Relationship with PPR

- Consider a neural network with one hidden layer — aka a single-layer perceptron.
- Looking at it from the viewpoint of a PPR, we can write

$$\begin{aligned} g_m(\boldsymbol{\omega}'_m \mathbf{X}) &= \beta_m \sigma(\alpha_{0m} + \boldsymbol{\alpha}'_m \mathbf{X}) \\ &= \beta_m \sigma(\alpha_{0m} + \|\boldsymbol{\alpha}_m\| \boldsymbol{\omega}'_m \mathbf{X}). \end{aligned} \quad (2)$$

- Here,  $\boldsymbol{\omega}_m = \boldsymbol{\alpha}_m / \|\boldsymbol{\alpha}_m\|$ .
- Another way to look at the relationship in (2) is

$$g_m(V_m) = \beta_m \sigma(\alpha_{0m} + \|\boldsymbol{\alpha}_m\| V_m).$$

9

## Comments

- If there is no hidden layer, a (classification) neural network would be equivalent to multinomial logistic regression.
- There can be more than one hidden layer, which can be thought of as allowing a sort of hierarchy — aka a multi-layer perceptron.
- As Hastie et al. (2009, Sec. 11.5.3) explain, it is generally best to scale (standardize) the input (predictors).
- There is model fitting, and a lot of subtleties, to consider; extensive details are given (or referenced) in Hastie et al. (2009).

10

## Comments contd.

- Like PPR, neural networks can work very well when prediction — and not modelling — is the goal.
- As Hastie et al. (2009) point out, neural networks (and PPR) generally work well in situations where there is high signal-to-noise ratio.
- In “competitions”, neural networks are often amongst the best learning methods.
- Let’s look at some neural network examples in R.